

VRANJAN_Week6.2a

April 26, 2021

```
[1]: # Course DSC 650 - Data Mining
     # Name - Vikas Ranjan
     # Assignment - Assignment 6.2a CIFAR10 small images classification dataset.
```

```
[2]: # Import packages
     from keras.datasets import cifar10
     from keras.utils import to_categorical
     from keras import models
     from keras import layers
     import matplotlib.pyplot as plt
     import matplotlib.pyplot as plt
     import pandas as pd
```

```
[3]: (x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

```
[4]: x_train.shape, y_train.shape
```

```
[4]: ((50000, 32, 32, 3), (50000, 1))
```

```
[5]: x_test.shape, y_test.shape
```

```
[5]: ((10000, 32, 32, 3), (10000, 1))
```

```
[6]: # Preprocess the data (these are NumPy arrays)
     x_train = x_train.astype("float32") / 255
     x_test = x_test.astype("float32") / 255

     y_train = to_categorical(y_train)
     y_test = to_categorical(y_test)

     # Reserve 10,000 samples for validation
     x_val = x_train[-10000:]
     y_val = y_train[-10000:]
     x_train = x_train[:-10000]
     y_train = y_train[:-10000]
```

```
[7]: x_val.shape, y_val.shape
```

```
[7]: ((10000, 32, 32, 3), (10000, 10))
```

```
[8]: #instantiate the model
model = models.Sequential()
model.add(layers.Conv2D(32, (3,3), activation='relu', input_shape=(32,32,3)))
model.add(layers.MaxPooling2D(2,2))
model.add(layers.Conv2D(64, (3,3), activation='relu'))
model.add(layers.MaxPooling2D(2,2))
model.add(layers.Conv2D(64, (3,3), activation='relu'))
model.add(layers.MaxPooling2D(2,2))
model.add(layers.Flatten())
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))

model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_2 (Conv2D)	(None, 4, 4, 64)	36928
max_pooling2d_2 (MaxPooling2D)	(None, 2, 2, 64)	0
flatten (Flatten)	(None, 256)	0
dense (Dense)	(None, 64)	16448
dense_1 (Dense)	(None, 10)	650

Total params: 73,418
Trainable params: 73,418
Non-trainable params: 0

```
[9]: model.compile(optimizer='rmsprop',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])
```


```
history = model.fit(x_train, y_train, epochs=100, validation_data=(x_val, y_val), verbose=0)
```

```
[10]: train_loss = history.history['loss']
      val_loss = history.history['val_loss']

      epochs = range(1, len(history.history['loss']) + 1)

      plt.plot(epochs, train_loss, 'bo', label='Training loss')
      plt.plot(epochs, val_loss, 'b', label='Validation loss')
      plt.title('Training and Validation Losses')
      plt.xlabel('Epochs')
      plt.ylabel('Loss')
      plt.legend()

      plt.show()
      plt.savefig('results/6_2a_lossplot.png')
```



<Figure size 432x288 with 0 Axes>

```
[11]: train_loss = history.history['accuracy']
      val_loss = history.history['val_accuracy']
```


```

epochs = range(1, len(history.history['accuracy']) + 1)

plt.plot(epochs, train_loss, 'bo', label='Training accuracy')
plt.plot(epochs, val_loss, 'b', label='Validation accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
plt.savefig('results/6_2a_accplot.png')

```



<Figure size 432x288 with 0 Axes>

```

[12]: #retrain the model and evaluate on test
(x_train, y_train), (x_test, y_test) = cifar10.load_data()

# Preprocess the data (these are NumPy arrays)
x_train = x_train.astype("float32") / 255
x_test = x_test.astype("float32") / 255

y_train = to_categorical(y_train)
y_test = to_categorical(y_test)

```

```

model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

history = model.fit(x_train, y_train, epochs=10)
results = model.evaluate(x_test, y_test)

```

```

Epoch 1/10
1563/1563 [=====] - 17s 10ms/step - loss: 0.9287 -
accuracy: 0.7639
Epoch 2/10
1563/1563 [=====] - 16s 10ms/step - loss: 0.7448 -
accuracy: 0.7836
Epoch 3/10
1563/1563 [=====] - 16s 10ms/step - loss: 0.6829 -
accuracy: 0.7947
Epoch 4/10
1563/1563 [=====] - 16s 10ms/step - loss: 0.6692 -
accuracy: 0.7935
Epoch 5/10
1563/1563 [=====] - 16s 11ms/step - loss: 0.6515 -
accuracy: 0.8005
Epoch 6/10
1563/1563 [=====] - 16s 11ms/step - loss: 0.6412 -
accuracy: 0.8008
Epoch 7/10
1563/1563 [=====] - 16s 10ms/step - loss: 0.6278 -
accuracy: 0.8024
Epoch 8/10
1563/1563 [=====] - 17s 11ms/step - loss: 0.6112 -
accuracy: 0.8042
Epoch 9/10
1563/1563 [=====] - 17s 11ms/step - loss: 0.6057 -
accuracy: 0.8090
Epoch 10/10
1563/1563 [=====] - 18s 11ms/step - loss: 0.5890 -
accuracy: 0.8118
313/313 [=====] - 2s 5ms/step - loss: 1.3643 -
accuracy: 0.6350

```

```
[13]: model.save('results/6_2a_model.h5')
```

```
[14]: prediction_results = model.predict(x_test)
```

```
[15]: #write metrics to file
with open('results/6_2a_metrics.txt', 'w') as f:
```

```
f.write('Training Loss: {}'.format(str(history.history['loss'])))  
f.write('\nTraining Accuracy: {}'.format(str(history.history['accuracy'])))  
f.write('\nTest Loss: {}'.format(results[0]))  
f.write('\nTest Accuracy: {}'.format(results[1]))
```

```
[16]: predictions = pd.DataFrame(prediction_results,   
    ↪ columns=['0', '1', '2', '3', '4', '5', '6', '7', '8', '9'])  
predictions.to_csv('results/6_2a_predictions.csv', index=False)
```

```
[ ]:
```