

VRANJAN_Week6.2b

April 26, 2021

```
[1]: # Course DSC 650 - Data Mining
# Name - Vikas Ranjan
# Assignment - Assignment 6.2b - Create a ConvNet model that classifies images_
→CIFAR10 small images classification dataset.
#
# This time includes dropout and data-augmentation.
→
```

```
[2]: from keras.datasets import cifar10
from keras.utils import to_categorical
from keras.preprocessing.image import ImageDataGenerator
from keras import models
from keras import layers
from keras import optimizers
import matplotlib.pyplot as plt
import pandas as pd
```

```
[3]: (x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

```
[4]: x_train.shape, y_train.shape
```

```
[4]: ((50000, 32, 32, 3), (50000, 1))
```

```
[5]: x_test.shape, y_test.shape
```

```
[5]: ((10000, 32, 32, 3), (10000, 1))
```

```
[6]: # Preprocess the data (these are NumPy arrays)
x_train = x_train.astype("float32")
x_test = x_test.astype("float32")

y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
```

```
[7]: # Reserve 10,000 samples for validation
x_val = x_train[-10000:]
y_val = y_train[-10000:]
x_train_2 = x_train[:-10000]
```

```
y_train_2 = y_train[:-10000]
```

```
[8]: train_datagen = ImageDataGenerator(rescale=1./255,
                                         rotation_range=40,
                                         width_shift_range=0.2,
                                         height_shift_range=0.2,
                                         shear_range=0.2,
                                         zoom_range=0.2,
                                         horizontal_flip=True)

test_datagen = ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow(x_train_2, y_train_2, batch_size=32)

validation_generator = train_datagen.flow(x_val, y_val, batch_size=32)
```

```
[9]: #instantiate the model
model = models.Sequential()
model.add(layers.Conv2D(32, (3,3), activation='relu', input_shape=(32,32,3)))
model.add(layers.MaxPooling2D(2,2))
model.add(layers.Conv2D(64, (3,3), activation='relu'))
model.add(layers.MaxPooling2D(2,2))
model.add(layers.Conv2D(64, (3,3), activation='relu'))
model.add(layers.MaxPooling2D(2,2))
model.add(layers.Flatten())
model.add(layers.Dropout(0.5))
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))

model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_2 (Conv2D)	(None, 4, 4, 64)	36928
max_pooling2d_2 (MaxPooling2D)	(None, 2, 2, 64)	0

flatten (Flatten)	(None, 256)	0

dropout (Dropout)	(None, 256)	0

dense (Dense)	(None, 64)	16448

dense_1 (Dense)	(None, 10)	650
=====		
Total params: 73,418		
Trainable params: 73,418		
Non-trainable params: 0		

```
[10]: model.compile(optimizer=optimizers.RMSprop(lr=1e-4),
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])
```


```
[11]: history = model.fit_generator(train_generator,
                                   steps_per_epoch=len(x_train_2) / 32,
                                   epochs=10,
                                   validation_data=validation_generator,
                                   validation_steps=len(x_val) / 32)
```

```
/opt/conda/lib/python3.8/site-
packages/tensorflow/python/keras/engine/training.py:1844: UserWarning:
`Model.fit_generator` is deprecated and will be removed in a future version.
Please use `Model.fit`, which supports generators.
  warnings.warn("`Model.fit_generator` is deprecated and "
```

```
Epoch 1/10
1250/1250 [=====] - 50s 40ms/step - loss: 2.2648 -
accuracy: 0.1389 - val_loss: 2.0724 - val_accuracy: 0.2226
Epoch 2/10
1250/1250 [=====] - 48s 39ms/step - loss: 2.0608 -
accuracy: 0.2241 - val_loss: 1.9252 - val_accuracy: 0.2807
Epoch 3/10
1250/1250 [=====] - 53s 42ms/step - loss: 1.9486 -
accuracy: 0.2648 - val_loss: 1.8692 - val_accuracy: 0.3027
Epoch 4/10
1250/1250 [=====] - 57s 45ms/step - loss: 1.8946 -
accuracy: 0.2900 - val_loss: 1.8183 - val_accuracy: 0.3214
Epoch 5/10
1250/1250 [=====] - 59s 47ms/step - loss: 1.8463 -
accuracy: 0.3073 - val_loss: 1.7770 - val_accuracy: 0.3503
Epoch 6/10
1250/1250 [=====] - 57s 46ms/step - loss: 1.7973 -
accuracy: 0.3280 - val_loss: 1.7426 - val_accuracy: 0.3615
Epoch 7/10
```

```
1250/1250 [=====] - 63s 50ms/step - loss: 1.7680 -  
accuracy: 0.3487 - val_loss: 1.7117 - val_accuracy: 0.3758  
Epoch 8/10  
1250/1250 [=====] - 58s 46ms/step - loss: 1.7363 -  
accuracy: 0.3590 - val_loss: 1.6827 - val_accuracy: 0.3841  
Epoch 9/10  
1250/1250 [=====] - 55s 44ms/step - loss: 1.7163 -  
accuracy: 0.3707 - val_loss: 1.6479 - val_accuracy: 0.4068  
Epoch 10/10  
1250/1250 [=====] - 60s 48ms/step - loss: 1.7057 -  
accuracy: 0.3734 - val_loss: 1.6321 - val_accuracy: 0.4098
```

```
[12]: train_loss = history.history['loss']  
val_loss = history.history['val_loss']  
  
epochs = range(1, len(history.history['loss']) + 1)  
  
plt.plot(epochs, train_loss, 'bo', label='Training loss')  
plt.plot(epochs, val_loss, 'b', label='Validation loss')  
plt.title('Training and Validation Losses')  
plt.xlabel('Epochs')  
plt.ylabel('Loss')  
plt.legend()  
  
plt.show()  
plt.savefig('results/6_2b_lossplot.png')
```




<Figure size 432x288 with 0 Axes>

```
[13]: train_loss = history.history['accuracy']
      val_loss = history.history['val_accuracy']

      epochs = range(1, len(history.history['accuracy']) + 1)

      plt.plot(epochs, train_loss, 'bo', label='Training accuracy')
      plt.plot(epochs, val_loss, 'b', label='Validation accuracy')
      plt.title('Training and Validation Accuracy')
      plt.xlabel('Epochs')
      plt.ylabel('Accuracy')
      plt.legend()

      plt.show()
      plt.savefig('results/6_2b_accplot.png')
```



<Figure size 432x288 with 0 Axes>

```
[14]: #retrain the model and evaluate on test
      train_generator = train_datagen.flow(x_train, y_train, batch_size=32)
```

```

model.compile(optimizer=optimizers.RMSprop(lr=1e-4),
              loss='categorical_crossentropy',
              metrics=['accuracy'])

#16 epochs chosen based on graphs above
history = model.fit_generator(train_generator,
                             steps_per_epoch=len(x_train) / 32,
                             epochs=16)
results = model.evaluate(x_test, y_test)

```

```

Epoch 1/16
1562/1562 [=====] - 55s 35ms/step - loss: 1.6707 -
accuracy: 0.3881
Epoch 2/16
1562/1562 [=====] - 58s 37ms/step - loss: 1.6496 -
accuracy: 0.4003
Epoch 3/16
1562/1562 [=====] - 62s 40ms/step - loss: 1.6299 -
accuracy: 0.4097
Epoch 4/16
1562/1562 [=====] - 53s 34ms/step - loss: 1.6135 -
accuracy: 0.4187
Epoch 5/16
1562/1562 [=====] - 53s 34ms/step - loss: 1.5988 -
accuracy: 0.4226
Epoch 6/16
1562/1562 [=====] - 55s 35ms/step - loss: 1.5733 -
accuracy: 0.4308
Epoch 7/16
1562/1562 [=====] - 62s 39ms/step - loss: 1.5712 -
accuracy: 0.4333
Epoch 8/16
1562/1562 [=====] - 62s 40ms/step - loss: 1.5525 -
accuracy: 0.4396
Epoch 9/16
1562/1562 [=====] - 62s 40ms/step - loss: 1.5411 -
accuracy: 0.4522
Epoch 10/16
1562/1562 [=====] - 62s 40ms/step - loss: 1.5293 -
accuracy: 0.4477
Epoch 11/16
1562/1562 [=====] - 60s 39ms/step - loss: 1.5218 -
accuracy: 0.4491
Epoch 12/16
1562/1562 [=====] - 58s 37ms/step - loss: 1.4960 -
accuracy: 0.4639

```

```

Epoch 13/16
1562/1562 [=====] - 64s 41ms/step - loss: 1.5001 -
accuracy: 0.4621
Epoch 14/16
1562/1562 [=====] - 65s 42ms/step - loss: 1.4826 -
accuracy: 0.4685
Epoch 15/16
1562/1562 [=====] - 64s 41ms/step - loss: 1.4774 -
accuracy: 0.4712
Epoch 16/16
1562/1562 [=====] - 53s 34ms/step - loss: 1.4679 -
accuracy: 0.4743
313/313 [=====] - 2s 5ms/step - loss: 178.5847 -
accuracy: 0.4082

```

```
[15]: model.save('results/6_2b_model.h5')
```

```
[16]: prediction_results = model.predict(x_test)
```

```
[17]: #write metrics to file
with open('results/6_2b_metrics.txt', 'w') as f:
    f.write('Training Loss: {}'.format(str(history.history['loss'])))
    f.write('\nTraining Accuracy: {}'.format(str(history.history['accuracy'])))
    f.write('\nTest Loss: {}'.format(results[0]))
    f.write('\nTest Accuracy: {}'.format(results[1]))
```

```
[18]: predictions = pd.DataFrame(prediction_results,
    ↪columns=['0','1','2','3','4','5','6','7','8','9'])
predictions.to_csv('results/6_2b_predictions.csv', index=False)
```